

# Timings of common tasks using the **data.table** package in R

Matthew Dowle

April 12, 2010

This document contains a series of tests, followed by a summary table of various timings and comparisons. Please go straight to the summary table first [<here>](#) in which each row has a link back to the test.

This document is reproducible. Simply run the .Rnw file yourself in your environment to confirm the results. Also see `?vignette`, which says that `edit(vignette("datatable-timings"))` will extract the code from this document so you can easily work with it.

The .Rnw included in the package has  $N=10,000,000$ . This is a small number so that 'R CMD build' completes in a reasonable time (about 5 minutes). We don't want the nightly builds on R-Forge and CRAN to slow down just to run long timing comparisons. We have increased this to  $N=100,000,000$  ourselves, and included the output on the [datatable homepage \(<link>\)](#).

## Contents

<b>1</b>	<b>Timing tests</b>	<b>1</b>
1.1	Extraction	1
1.2	Grouping	2
<b>2</b>	<b>Summary table</b>	<b>2</b>

## 1 Timing tests

### 1.1 Extraction

This is a repeat of the test in section 1 of the Introduction vignette. The syntax is explained there. This demonstrates the large difference in speed between vector scans and binary search. Therefore, please avoid using `==` in the `i` expression.

```
> n = ceiling(1e7/26^2) # 10 million rows
> DT = data.table(x=rep(LETTERS,each=26*n),
+                y=rep(letters,each=n),
+                v=rnorm(n*26^2),
+                key="x,y")
> tables()

      NAME      NROW  MB COLS  KEY
[1,] DT    10,000,068 153  x,y,v x,y
Total: 153MB

> tt=system.time(ans1 <- DT[DT$x=="R" & DT$y=="h",]); tt
      user system elapsed
4.284    1.168    5.839

> ss=system.time(ans2 <- DT[J("R","h"),mult="all"]); ss
      user system elapsed
0.016    0.000    0.017

> if(!identical(ans1,ans2)) stop("Test 1 not identical")
```

## 1.2 Grouping

Repeat from Introduction.

## 2 Summary table

```
> tt[3]
```

```
elapsed  
5.839
```

```
> ss[3]
```

```
elapsed  
0.017
```