FAQs about the **data.table** package in R

Matthew Dowle

February 15, 2011
(A later revision may be available on the homepage)

The first section, Beginner FAQs, is intended to be read in order from start to finish. It may be read before reading the Introduction vignette.

Contents

1 Beginner FAQs

1.1 Why does `DT[,5]` return `5` ?

Because, by default, unlike a `data.frame`, the 2nd argument is an *expression* which is evaluated within the scope of `DT`. 5 evaluates to 5. It is generally bad practice to refer to columns by number rather than name. If someone else comes along and reads your code later, they may have to hunt around to find out which column is number 5. Furthermore, if you or someone else changes the column ordering of `DT` higher up in your R program, you might get bugs if you forget to change all the places in your code which refer to column number 5.

Say column 5 is called "region", just do `DT[,region]` instead. Notice there are no quotes around the column name. This is what we mean by j being evaluated within the scope of the `data.table`. That scope consists of an environment where the column names are variables.

You can write *any* R expression in the j e.g. `DT[,colA*colB/2]`. Further, j may be a `list()` of many R expressions, including calls to any R package e.g. `DT[,fitdistr(d1-d1,"normal")]`.

Having said this, there are some circumstances where referring to a column by number is ok, such as a sequence of columns. In these situations just do `DT[,5:10,with=FALSE]` or `DT[,c(1,4,10),with=FALSE]`. See `?data.table` for an explanation of the `with` argument.

Note that `with()` has been a base function for a long time. That's one reason we say `data.table` builds upon base functionality. There is little new here really, `data.table` is just making use of `with()` and building it into the syntax.

1.2 Why does `DT[,"region"]` return `"region"`?

See answer to 1.1 above. Try `DT[,region]` instead. Or `DT[,"region",with=FALSE]`.

1.3 Why does `DT[,region]` return a vector? I'd like a 1-column `data.table`. There is no drop argument like I'm used to in `data.frame`.

Try `DT[,list(region)]` instead.

1.4 Why does `DT[,x,y,z]` not work? I wanted the 3 columns x,y and z.

The `j` expression is the 2nd argument. The correct way to do this is `DT[,list(x,y,z)]`.

1.5 I assigned a variable `mycol="x"` but then `DT[,mycol]` returns `"x"`. How do I get it to look up the column name contained in the `mycol` variable?

This is what we mean when we say the `j` expression 'sees' objects in the calling scope. The variable `mycol` does not exist as a column name of DT so R then looked in the calling scope and found `mycol` there, and returned its value `"x"`. This is correct behaviour. Had `mycol` been a column name, then that column's data would have been returned. What you probably meant was `DT[,mycol,with=FALSE]`, which will return the `x` column's data as you wanted. Alternatively, since a `data.table` is just a list, you can do `DT[["x"]]` or `DT[[mycol]]`.

1.6 Ok but I don't know the expressions in advance. How do I programatically pass them in?

To create expressions use the `quote()` function. We refer to these as *quote()-ed* expressions to save confusion with the double quotes used to create a character vector such as `c("x")`. The simplest quote()-ed expression is just a column name on it's own :
```
q = quote(x)
DT[,eval(q)] # returns the column x as a vector
q = quote(list(x))
DT[,eval(q)] # returns the column x as a 1-column data.table
```
Since these are *expressions*, we are not restricted to column names only :
```
q = quote(mean(x))
DT[,eval(q)] # identical to DT[,mean(x)]
q = quote(list(x,sd(y),mean(y*z)))
DT[,eval(q)] # identical to DT[,list(x,sd(y),mean(y*z))]
```
If it's just simply a vector of column names it may be simpler to pass a character vector to `j` and use `with=FALSE`.

1.7 This is really hard. What's the point ?

`j` doesn't have to be just column names. You can put any R expression of column names directly as the `j`, e.g., `DT[,mean(x*y/z)]`. The same applies to `i`. You have been used to `i` being row numbers or row names only. It's nice to simply write `DT[x>1000, sum(y*z)]`. What does that mean? It just runs the `j` expression on the set of rows where the `i` expression is true. `i` can be any expression of column names that evaluates to logical. You don't even need to return data, e.g., `DT[x>1000, plot(y,z)]`. When we get to compound table joins we will see how `i` and `j` can themselves be other `data.table` queries. We are going to stretch `i` and `j` much further than this, but to get there we need you on board first with FAQs 1.1-1.6.

1.8 OK, I'm starting to see what `data.table` is about, but why didn't you enhance `data.frame` in **R**? Why does it have to be a new package?

As FAQ 1.1 highlights, j in `[.data.table` is fundamentally different from j in `[.data.frame`. Even something as simple as `DF[,1]` would break existing code in many packages and user code. This is by design, and we want it to work this way for more complicated syntax to work.

However, we *have* proposed enhancements to R wherever possible. Some of these have been accepted and included. We intend to continue attempts to contribute to R itself whenever an opportunity arises.

Since `data.table` *inherits* from `data.frame` it can be passed to any package and that package can use `[.data.frame` syntax on `data.table` and it just works. See `?data.table`.

1.9 Why are the defaults the way they are? Why does it work the way it does?

The simple answer is because the author designed it for his own use, and he wanted it that way. He finds it a more natural, faster way to write code, which also executes more quickly.

1.10 Isn't this already done by `with()` and `subset()` in base?

Some of the features discussed so far are, yes. The package builds upon base functionality. It does the same sorts of things but with less code required, and executes many times faster if used correctly.

1.11 Why does `X[Y]` return all the columns from `Y` too? Shouldn't it return a subset of `X`?

This was changed in v1.5.3. `X[Y]` now includes `Y`'s non-join columns. We refer to this feature as *join inherited scope* because not only are `X` columns available to the j expression, so are `Y` columns. The downside is that `X[Y]` is less efficient since every item of `Y`'s non-join columns are duplicated to match the (likely large) number of rows in `X` that match. We therefore strongly encourage `X[Y,j]` instead of `X[Y]`. See next FAQ.

1.12 What is the difference between `X[Y]` and `merge(X,Y)`?

`X[Y]` is a join, looking up `X`'s rows using `Y` (or `Y`'s key if it has one) as an index.
`Y[X]` is a join, looking up `Y`'s rows using `X` (or `X`'s key if it has one) as an index.
`merge(X,Y)` does both ways at the same time. The number of rows of `X[Y]` and `Y[X]` usually differ; the number of rows returned by `merge(X,Y)` and `merge(Y,X)` is the same.

BUT that misses the main point. Most tasks require something to be done on the data after a join or merge. Why merge all the columns of data, only to use a small subset of them afterwards? You may suggest `merge(X[,ColsNeeded1],Y[,ColsNeeded2])`, but that takes copies of the subsets of data, and it requires the programmer to work out which columns are needed. `X[Y,j]` in data.table does all that in one step for you. When you write `X[Y,sum(foo*bar)]`, data.table automatically inspects the j expression to see which columns it uses. It will only subset those columns only; the others are ignored. Memory is only created for the columns the j uses, and `Y` columns enjoy standard R recycling rules within the context of each group. Let's say `foo` is in `X`, and `bar` is in `Y` (along with 20 other columns in `Y`). Isn't `X[Y,sum(foo*bar)]` quicker to program and quicker to run than a `merge` followed by a `[`?

1.13 Anything else about `X[Y,sum(foo*bar)]`?

Remember that j (in this example `sum(foo*bar)`) is run for each *group* of `X` that each row of `Y` matches to. This feature is *grouping by i* or *by without by*.

1.14 That's nice but what if I really do want to evaluate `j` for all rows once, not by group?

If you really want `j` to run once for the whole subset of X then try `X[Y][,sum(foo*bar)]`. If that needs to be efficient (recall that `X[Y]` joins all columns) then you will have to work a little harder since this is outside the common use-case: `X[Y,list(foo,bar)][,sum(foo*bar)]`.

2 General syntax

2.1 How can I avoid writing a really long `j` expression? You've said I should use the column names, but I've got a lot of them.

One way is to use `data.table` in the `data.frame` way by setting `with=FALSE` and using `sapply` or `lapply` on the columns. A second way is to *flatten* the table. Example required. Finally, when grouping, the `j` expression can use column names as variables, as you know, but it can also use a reserved symbol `.SD` which refers to the subset of the `data.table` for each group (excluding the grouping columns). So to sum up all your columns it's just `DT[,lapply(.SD,sum),by=grp]`. It might seem tricky, but it's fast to write and fast to run. Notice you don't have to create an anonymous `function()`. See the timing vignette and comparison to other methods. The `.SD` object is efficiently implemented internally and more efficient than passing an argument to a function. Please don't do this though : `DT[,.SD["sales",with=FALSE],by=grp]`. That works but is very inefficient and inelegant. This is what was intended: `DT[,sum(sales),by=grp]` and could be 100's of times faster.

2.2 Why is the default for mult now `"all"`?

In v1.5.3 the default was changed to `"all"`. When `i` (or `i`'s key if it has one) has fewer columns than `x`'s key, `mult` was already set to `"all"` automatically. Changing the default makes this clearer and easier for users as it came up quite often.

In versions up to v1.3, `"all"` was slower. Internally, `"all"` was implemented by joining using `"first"`, then again from scratch using `"last"`, after which a diff between them was performed to work out the span of the matches in `x` for each row in `i`. Most often we join to single rows, though, where `"first"`,`"last"` and `"all"` return the same result. We preferred maximum performance for the majority of situations so the default chosen was `"first"`. When working with a non-unique key (generally a single column containing a grouping variable), `DT["A"]` returned the first row of that group so `DT["A",mult="all"]` was needed to return all the rows in that group.

In v1.4 the binary search in C was changed to branch at the deepest level to find first and last. That branch will likely occur within the same final pages of RAM so there should no longer be a speed disadvantage in defaulting `mult` to `"all"`. We warned that the default might change, and made the change in v1.5.3.

A future version of `data.table` may allow a distinction between a key and a *unique key*. Internally `mult="all"` would perform more like `mult="first"` when all `x`'s key columns were joined to and `x`'s key was a unique key. `data.table` would need checks on insert and update to make sure a unique key is maintained. An advantage of specifying a unique key would be that `data.table` would ensure no duplicates could be inserted, in addition to performance.

2.3 I'm using `c()` in the `j` and getting strange results.

This is a common source of confusion. In `data.frame` you are used to, for example:

```
> DF = data.frame(x=1:3,y=4:6,z=7:9)
> DF

  x y z
1 1 4 7
2 2 5 8
3 3 6 9
```

```
> DF[,c("y","z")]

  y z
1 4 7
2 5 8
3 6 9
```

which returns the two columns. In `data.table` you know you can use the column names directly and might try :

```
> DT = data.table(DF)
> DT[,c(y,z)]

[1] 4 5 6 7 8 9
```

but this returns one vector. Remember that the j expression is evaluated within the environment of `DT`, and `c()` returns a vector. If 2 or more columns are required, use list() instead:

```
> DT[,list(y,z)]

     y z
[1,] 4 7
[2,] 5 8
[3,] 6 9
```

`c()` can be useful in a `data.table` too, but its behaviour is different from how it works in `[.data.frame`.

2.4 I have built up a complex table with many columns. I want to use it as a template for a new table; i.e., create a new table with no rows, but with the column names and types copied from my table. Can I do that easily?

Yes. If your complex table is called `DT`, try `NEWDT = DT[0]`.

2.5 Is a NULL `data.table` the same as `DT[0]`?

No, despite the print method indicating otherwise. Strictly speaking it's not possible to have `is.null(data.table(NULL))` return FALSE. This FAQ may be revisited in future.

2.6 Why has the `DT()` alias been removed?

`DT` was introduced originally as a wrapper for a list of j expressions. Since `DT` was an alias for `data.table`, this was a convenient way to take care of silent recycling in cases where each item of the j list evaluated to different lengths. The alias was one reason grouping was slow, though. As of v1.3, `list()` should be passed instead to the j argument. `list()` is a primitive and is much faster, especially when there are many groups. Internally, this was a nontrivial change. Vector recycling is done internally, along with several other speed enhancements for grouping.

2.7 But my code uses j=DT(...) and it works. The previous FAQ says that DT() has been removed.

Then you are using a version prior to 1.5.3. Prior to 1.5.3 `[.data.table` detected use of `DT()` in the j and automatically replaced it with a call to `list()`. This was to help the transition for existing users.

2.8 What are the scoping rules for `j` expressions?

Think of the subset as an environment where all the column names are variables. When a variable `foo` is used in the `j` of a query such as `X[Y,sum(foo)]`, `foo` is looked for in the following order :

1. The scope of `X`'s subset i.e. `X`'s column names.

2. The scope of each row of `Y` i.e. `Y`'s column names (*join inherited scope*)

3. The scope of the calling frame e.g. the line that appears before the `data.table` query.

4. Exercise for reader: does it then ripple up the calling frames, or go straight to `globalenv()`?

5. The global environment

This is *lexical scoping* as explained in R FAQ 3.3.1. The environment that the function was created is not relevant, though, because there is *no function*. No *function* (anonymous nor named) is passed to the `j`. Instead, an anonymous *body* is passed to the `j`, e.g.

```
> DT = data.table(x=rep(c("a","b"),c(2,3)),y=1:5)
> DT

     x y
[1,] a 1
[2,] a 2
[3,] b 3
[4,] b 4
[5,] b 5

> DT[,{z=sum(y);z+3},by=x]

     x V1
[1,] a  6
[2,] b 15
```

Some programming languages call this a `lambda`.

2.9 Can I trace the `j` expression as it runs through the groups?

Try something like this:

```
> DT[,{
+   cat("Objects:",paste(objects(),collapse=","),"\n")
+   cat("Trace: x=",as.character(x)," y=",y,"\n")
+   sum(y)
+ },by=x]

Objects: x,y
Trace: x= a a  y= 1 2
Objects: x,y
Trace: x= b b b  y= 3 4 5
     x V1
[1,] a  3
[2,] b 12
```

2.10 Inside each group, why is the group variable a long vector containing the same value repeated?

This is correct. We saw that in the previous FAQ. x was `"a"` repeated twice in the first group, and `"b"` repeated three times in the second group. When you group, think of the data being split up. Sometimes we want to use the value of the group in the expression, though. In that case, we just use the first value.

```
> DT[,list(g=1,h=2,i=3,j=4,repeatgroupname=x,sum(y)),by=x]      # not intended

     x g h i j repeatgroupname V6
[1,] a 1 2 3 4               a  3
[2,] a 1 2 3 4               a  3
[3,] b 1 2 3 4               b 12
[4,] b 1 2 3 4               b 12
[5,] b 1 2 3 4               b 12

> DT[,list(g=1,h=2,i=3,j=4,repeatgroupname=x[1],sum(y)),by=x]  # intended

     x g h i j repeatgroupname V6
[1,] a 1 2 3 4               a  3
[2,] b 1 2 3 4               b 12
```

In the first attempt, the aggregate sum(y) was recycled to match the length of x. Recycling can be useful, but wasn't intended here.

2.11   Only the first 10 rows are printed, how do I print more?

Try print(DT,nrows=Inf) to print all rows, or set nrows to the number of rows you require.

2.12   With an X[Y] join, what if X contains a column called "Y"?

When i is a single name such as Y it is evaluated in the calling frame. In all other cases such as calls to J() or other expressions, i is evaluated within the scope of X. This facilitates easy *self joins* such as X[J(unique(colA)),mult="first"].

2.13   X[Z[Y]] is failing because X contains a column "Y". I'd like it to use the table Y in calling scope.

The Z[Y] part is not a single name so that is evaluated within the frame of X and the problem occurs. Try tmp=Z[Y];X[tmp]. This is robust to X containing a column "tmp" because tmp is a single name. If you often encounter conflics of this type, one simple solution may be to name all tables in uppercase and all column names in lowercase, or some similar scheme.

2.14   Can you explain further why data.table is inspired by A[B] syntax in base?

Consider A[B] syntax using an example matrix A.

```
> A = matrix(1:12,nrow=4)
> A

     [,1] [,2] [,3]
[1,]    1    5    9
[2,]    2    6   10
[3,]    3    7   11
[4,]    4    8   12
```

To obtain cells (1,2)=5 and (3,3)=11 many users (we believe) may try this first :

```
> A[c(1,3),c(2,3)]

     [,1] [,2]
[1,]    5    9
[2,]    7   11
```

That returns the union of those rows and columns, though. To reference the cells, a 2-column matrix is required. See what `?Extract` says :

> When indexing arrays by [a single argument i can be a matrix with as many columns as there are dimensions of x; the result is then a vector with elements corresponding to the sets of indices in each row of i.

Let's try again.

```
> B = cbind(c(1,3),c(2,3))
> B

     [,1] [,2]
[1,]    1    2
[2,]    3    3

> A[B]

[1]  5 11
```

A matrix is a 2-dimension structure with row names and column names. Can we do the same with names?

```
> rownames(A) = letters[1:4]
> colnames(A) = LETTERS[1:3]
> A

  A B  C
a 1 5  9
b 2 6 10
c 3 7 11
d 4 8 12

> B = cbind(c("a","c"),c("B","C"))
> A[B]

[1]  5 11
```

So, yes we can. Can we do the same with `data.frame`?

```
> A = data.frame(A=1:4,B=letters[11:14],C=pi*1:4)
> rownames(A) = letters[1:4]
> A

  A B         C
a 1 k  3.141593
b 2 l  6.283185
c 3 m  9.424778
d 4 n 12.566371

> B

     [,1] [,2]
[1,] "a"  "B"
[2,] "c"  "C"

> A[B]

[1] "k"        " 9.424778"
```

But, notice that the result was coerced to character. R coerced `A` to matrix first so that the syntax could work, but the result isn't ideal. Let's try making `B` a `data.frame`.

```
> B = data.frame(c("a","c"),c("B","C"))
> cat(try(A[B],silent=TRUE))

Error in `[.default`(A, B) : invalid subscript type 'list'
```

So we can't subset a `data.frame` by a `data.frame` in base R. What if we want row names and column names that aren't character but integer or float? What if we want more than 2 dimensions of mixed types?

Matrices, especially sparse matrices, are often stored in a 3 column tuple: (i,j,value). This can be thought of as a key-value pair where `i` and `j` form a 2-column key. If we have more than one value, perhaps of different types it might look like (i,j,val1,val2,val3,...). This looks very much like a `data.frame`. Hence `data.table` extends `data.frame` so that a `data.frame` X can be subset by a data.frame Y, leading to the `X[Y]` syntax.

2.15 Can base be changed to do this then, rather than a new package?

Perhaps it can, we don't know. Perhaps if people like it and there is enough demand. Bear in mind that `data.frame` is used *everywhere* and so it is very difficult to make *any* changes to it. Is it really too much trouble to use `data.table`?

2.16 I've heard that `data.table` syntax is analogous to SQL?

Yes :

- `i` <==> where

- `j` <==> select

- `by` <==> group by

- `i` <==> order by (in compound syntax)

- `i` <==> having (in compound syntax)

- `nomatch=NA` <==> outer join

- `nomatch=0` <==> inner join

- `mult="first"|"last"` <==> N/A because SQL is inherently unordered

The general form is :
```
            DT[where,select,group by][having][order by][ ]...[ ]
```

A key advantage of column vectors in R is that they are *ordered*, unlike SQL. We can use ordered functions in `data.table` queries, such as `diff()`, and we can use *any* R function from any package, not just the functions that are defined in SQL. A disadvantage is that R objects must fit in memory, but with several R packages such as ff, bigmemory and indexing, this is changing.

2.17 What are the smaller syntax differences between `data.frame` and `data.table`?

- `DT[3]` refers to the 3rd row, but `DF[3]` refers to the 3rd column

- `DT[3,]` == `DT[3]`, but `DF[,3]` == `DF[3]` (somewhat confusingly)

- For this reason we say the comma is *optional* in `DT`, but not optional in `DF`

- `DT[[3]]` == `DF[3]` == `DF[[3]]`

- `DT[i,]` where `i` is a single integer returns a single row, just like `DF[i,]`, but unlike a matrix single row subset which returns a vector.

- `DT[,j,with=FALSE]` where j is a single integer returns a one column `data.table`, unlike `DF[,j]` which returns a vector by default

- There is no `drop` in DT. Instead, `DT[,1][[1]] == DF[,1]`.

- `DT[,colname] == DF[,"colname"]`

- `DT[,list(colname)] == DF[,"colname",drop=FALSE]`

When you pass a `data.table` to a `data.table`-unaware package, that package it not concerned with any of these differences. It just works.

2.18   I'm using `j` for it's side effect only, but I'm still getting data returned. How do I stop that?

In this case j can be wrapped with `invisible()` eg: `DT[,invisible(hist(colB)),by=colA]`. Here `hist()` returns the breakpoints in addition to plotting to the graphics device.

2.19   Why does `[.data.table` now have a `drop` argument from v1.5?

So that `data.table` can inherit from `data.frame` without using `....` If we used `...` then invalid argument names would not be caught and the convenience of test 147 would be lost.

The `drop` argument is never used by `[.data.table`. It is a placeholder for non `data.table` aware packages when they use the `[.data.frame` syntax directly on a `data.table`.

3   Questions relating to compute time

3.1   I have 20 columns in `data.table x`. Why is an expression of one column so quick?

Several reasons:

- Only that column is grouped, the other 19 are ignored because `data.table` inspects the j expression and realises it doesn't use the other columns.

- One memory allocation is made for the largest group only, then that memory is re-used for the other groups, there is very little garbage to collect.

- R is an in memory column store i.e. the columns are contiguous in RAM. Page fetches from RAM into L2 cache are minimised.

3.2   I don't have a key on a large table, but grouping is still really quick. Why is that?

`data.table` uses radix sorting. This is significantly faster than other sort algorithms. Radix is specifically for integers only, see `?base::sort.list(x,method="radix")`.

This is also one reason why `setkey()` is quick.

When no key is set, or we group in a different order from that of the key, we call it an *ad hoc by*.

3.3   Why is grouping by columns in the key faster than an ad hoc by?

Because each group is contiguous in RAM, thereby minimising page fetches, and memory can be copied in bulk (memcpy in C) rather than looping in C.

4 Error messages

4.1 Could not find function "DT"

See FAQ 2.6 and FAQ 2.7.

4.2 unused argument(s) (MySum = sum(v))

This error is generated by DT[,MySum=sum(v)]. DT[,list(MySum=sum(v))] was intended, or DT[,j=list(MySum=sum(v))].

5 General questions about the package

5.1 v1.3 appears to be missing from the CRAN archive ?

That is correct. v1.3 was available on R-Forge only. There were several large changes internally and these took some time to test in development.

5.2 Is data.table compatible with S-plus ?

Not currently.

- A few core parts of the package are written in C and use internal R functions and R structures.

- The package uses lexical scoping which is one of the differences between R and S-plus explained by R FAQ 3.3.1.

5.3 Is it available for Linux, Mac and Windows?

Yes, for both 32-bit and 64-bit on all platforms. Thanks to CRAN and R-Forge. There are no special or OS-specific libraries used.

5.4 I think it's great. What can I do?

Please send suggestions, bug reports and enhancement requests to datatable-help@lists.r-forge.r-project.org. This helps make the package better. The list is public and archived.

Please do vote for the package at http://crantastic.org/packages/data-table. This helps encourage the developers. If you have time to write a comment too, that can help others in the community. Just simply clicking that you use the package, though, is much appreciated.

You can join the project and change the code and/or documentation yourself.

5.5 I think it's not great. How do I warn others about my experience?

Please put your vote and comments on http://crantastic.org/packages/data-table. Please make it constructive so we have a chance to improve.

5.6 I have a question. I know the posting guide tells me to contact the maintainer (not r-help), but that's just two people. Is there a larger group of people I can ask?

Yes. You can post to datatable-help@lists.r-forge.r-project.org. It's like r-help, but just for this package. Feel free to answer questions there, too. maintainer("<package name>") returns the maintainer for any package and is where you should ask first (see r-help posting guide).

5.7 Where are the datatable-help archives?

The homepage contains links to the archives in 4 different formats.

5.8 I'd prefer not to contact datatable-help, can I mail just one or two people privately?

Sure. You're more likely to get a faster answer from datatable-help, though.

5.9 I have created a package that depends on `data.table`. How do I ensure my package is datatable-aware so that inheritance from data.frame still works?

You don't need to do anything special. As long as your package imports or depends on data.table, your package is detected as datatable-aware.

5.10 Why is this FAQ in pdf format? Can it moved to a website?

This FAQ (and the intro and timings) are *vignettes* written using Sweave. The benefits of Sweave include the following:

- We include R code in the vignettes. This code is *actually run* when the file is created, not copy and pasted.

- This document is *reproducible*. Grab the .Rnw and you can run it yourself.

- CRAN checks the package (including running vignettes) every night on Linux, Mac and Windows, both 32bit and 64bit. Results are posted to `http://cran.r-project.org/web/checks/check_results_data.table.html`. Included there are results from r-devel (i.e. not yet released R). That serves as a very useful early warning system for any potential future issues as R itself develops.

- This file is bound into each version of the package. The package is not accepted on CRAN unless this file passes checks. Each version of the package will have its own FAQ file which will be relevant for that version. Contrast this to a single website, which can be ambiguous if the answer depends on the version.

- You can open it offline, from your R prompt using `vignette()`.

- You can extract the code from the document and play with it using `edit(vignette("datatable-faq"))` or `edit(vignette("datatable-timings"))`.

- It prints out easily.

- It's quicker and easier for us to write and maintain the FAQ in .Rnw form.

Having said all that, a wiki format may be quicker and easier for users to contribute documentation and examples. Therefore a wiki has now been created; see link on the homepage.